



POLITECNICO
MILANO 1863

AUDIO EFFECTS CLASSIFICATION

GRUPPO 10

MEMBERS' PERSONAL CODES:

10635718 - 10743616 - 10776190 - 10787525

[Link](#) to GitHub repository

Abstract

The aim of this project was to implement a classifier able to identify the audio effect used in recordings of electric guitar and bass. In particular, we developed a system able to distinguish with high accuracy monophonic sounds produced by an electric guitar when **tremolo**, **distortion** or **no effects** are applied. The same model was applied on electric bass guitar's recordings, showing good accuracy as well, though less than the latter case. Noticeably, we exploited the open source Python library *sklearn* in order to choose the best model according to the features we used, and developed a *custom feature* to increase the accuracy of our model when selecting only 3 features.

Contents

Contents	i
List of Figures	iii
List of Tables	iii
1 Background	1
1.1 Database description	1
1.2 Features	1
1.2.1 Implementation	1
1.2.2 Best features selection	2
1.3 Model	2
1.3.1 Best model selection	2
2 Implementation	3
2.1 Data Gathering	3
2.2 Preprocessing	3
2.2.1 Amplitude Normalization	3
2.2.2 Length and Sampling frequency	4
2.3 Features Computation	4
2.4 Train-Test Splitting	4
2.5 Features Analysis	4
2.5.1 Features Normalization	4

2.5.2	Best Features Selection	5
2.6	Model Selection and Validation	5
2.6.1	All-features model	5
2.6.2	Selected-features model	6
2.6.3	Selected-features model - excluding our <i>custom feature</i>	7
3	Results And Discussions	10
3.1	Overall Conclusions On Accuracy Of Our Model	10
	Bibliography	11
	A1 List of features	12
	A2 List of models	14

List of Figures

2.1	Features plot for all the samples(note: most easily distinguishable is "Distortion" class, while "NoFx" and "Tremolo" exhibit quite similar properties)	5
2.2	Best all-features model confusion matrix	6
2.3	Best selected-features model confusion matrix when we consider our <i>custom feature</i>	7
2.4	Best selected-features model confusion matrix when we do not consider our <i>custom feature</i>	8
2.5	Features plot with the 3 k-best selected features	8
2.6	Scatter plot with the 3 k-best selected features	9

List of Tables

2.1	Table of scores for all-features models	6
2.2	Table of scores for selected-features models	7

1 | Background

1.1 Database description

The IDMT-SMT-Audio-Effects database is a large database from *Fraunhofer Institute for Digital Media Technology IDMT*. The database consists of .WAV files of monophonic and polyphonic sounds. It contains recordings of 2 different electric guitars and 2 different electric bass guitars, each with two different pick-up settings and up to three different plucking styles (finger plucked - hard, finger plucked - soft, picked) [1].

Among the 11 included audio effects, we focused, as requested, on the three classes **NoFx**, **Tremolo** and **Distortion**. In order to deal with the fact that for both the electric guitar and the electric bass guitar the NoFx class was underrepresented, we used a special parameter on *sklearn.svm.SVC* to weight the three classes differently, so that we did not need to augment the database by hand. The parameter is indeed `class_weight:balanced`, the “balanced” mode uses the values of *y* to automatically adjust weights inversely proportional to class frequencies in the input data [2].

1.2 Features

1.2.1 Implementation

To implement the features, most fundamental it has been the open source Python library *librosa*. For each of the following short time features, we have computed:

mean, minimum, maximum and standard deviation. Additionally, for the Mel-scaled Spectrogram, MFCC, and Spectral Contrast values we have added the delta features - local estimate of the derivative of the input data. Finally, we defined a *textitcustom* feature to better discriminate the "Tremolo class. We ended up with 365 values characterising each of the recordings[3]. See Appendix A1 for a list of the features we used, included our *custom feature*.

1.2.2 Best features selection

In order to identify the most useful features for our classification problem we imported from *sklearn.feature_selection* the *SelectKBest* utility which allows to select features according to the k highest scores.

1.3 Model

1.3.1 Best model selection

We exploited *GridSearchCV* from *sklearn.model_selection* to perform **hyperparameter tuning** in order to determine the optimal values for a given model. We pass predefined values for hyperparameters to the *GridSearchCV* function. To properly use the function, we defined a dictionary in which we mention a particular hyperparameter along with the values it can take. *GridSearchCV* tries all the combinations of the values passed in the dictionary and evaluates the model for each combination using the *Cross-Validation* method. Hence, we get accuracy/loss ratio for every combination of hyperparameters and we can choose the one with the best performance [4]. We used this strategy to identify the best classifier among a collection of 5 classifiers already implemented in *sklearn* library and define, at the same time, the best hyperparameters. See Appendix A2 for a list of the models we taken into account.

2 | Implementation

2.1 Data Gathering

As mentioned above, we started from monophonic electric guitar recordings. The recordings are actually divided into two folders, thus, for simplicity, we moved all the recordings of our interest (i.e., "NoFx", "Tremolo", "Distortion") in a dedicated directory. We gathered all the recordings in a list and collected the respective labels in a external .npy file. The overall number of recordings was 4368, 1872 "Distortion" recordings, 624 "NoFx" recordings, 1872 "Tremolo" recordings. This data-set is noticeably unbalanced.

2.2 Preprocessing

2.2.1 Amplitude Normalization

A plot of the minimum and maximum value in amplitude of each recording we noticed that the three classes are clearly distinguishable from their magnitude. To prevent the intensity of the sound to influence our model, we normalized the amplitude of our recordings from -1 to 1. As expected, this has shown to slight decrease the accuracy of our model.

2.2.2 Length and Sampling frequency

A quick check on the length of our recordings showed that they have been already preprocessed to have all the length of 88201 samples. They also have been recorded at the same sampling frequency of 44.1 kHz.

2.3 Features Computation

To properly compute the features we choose window parameters for windowing the signals according to the **Constant Overlapp-And-Add** condition. Thus, we selected a *Hamming* window of 1024 recordings, with hop size equal to half of the window length. We computed the features using the *librosa* library. For what concern *MFCC* and *Mel-Scaled Spectrogram*, we used a number of coefficient respectively equal to 21 and 40, as commonly reported in literature. The array of features was stored in a .npz file.

2.4 Train-Test Splitting

To randomly split the data-set in training-set and test-set we used *train_test_split* from *sklearn.model_selection*. This utility maintains unchanged the relative dimensions of each class and shuffles the recordings.

2.5 Features Analysis

2.5.1 Features Normalization

Normalization of the entire features array was done through the *minmax_scale* function from *sklearn.preprocessing* - the results were plot in figure 2.1. We used the *MinMaxScaler* object from the same package to normalize the training set and the test set, accordingly to the range defined by minimum and maximum values of our training set.




Figure 2.1: Features plot for all the samples (note: most easily distinguishable is "Distortion" class, while "NoFX" and "Tremolo" exhibit quite similar properties)

2.5.2 Best Features Selection

Using *SelectKBest* we found out the **3 most important features** to be the *standard deviation of Spectral Roll-off*, the *mean value of Spectral Flatness* and our *custom feature*. If had not taken into account our *custom feature*, the third selected feature would have been the *standard deviation of Spectral Centroid*.

2.6 Model Selection and Validation

2.6.1 All-features model

Best Model Selection

Considering a model which uses the entire set of features we defined, the best model turned out to be the *SVM*. The *GridSearchCV* function allowed us to find the best hyperparameters for our case. The worse model was instead the *KNN*.

	model	best_score	best_params
0	SVM	0.998569	{'C': 100, 'gamma': 0.05, 'kernel': 'rbf'}
1	Decision_Tree	0.997424	{'max_depth': 15}
2	Logistic_Regression	0.997138	{'C': 1}
3	Random_Forest	0.993989	{'max_depth': 15, 'n_estimators': 20}
4	KNN	0.979394	{'metric': 'minkowski', 'n_neighbors': 11}

Table 2.1: Table of scores for all-features models

Validation

Feeding our *SVM* with the test data-set we eventually got an incredibly high accuracy level, equal to 0.9965675057208238. Computing and plotting the *confusion matrix* from *sklearn.metrics* package, we could see that all the 3 classes are almost always identified.




Figure 2.2: Best all-features model confusion matrix

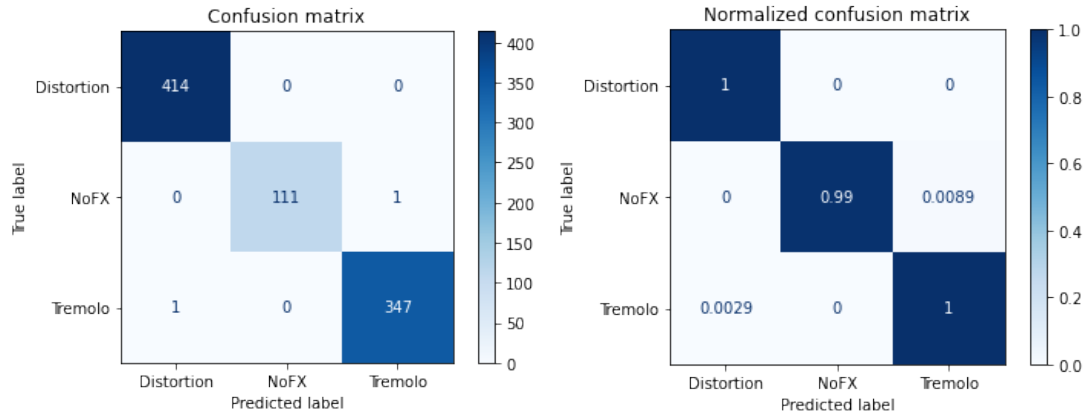
2.6.2 Selected-features model

Best Model Selection

Using only the 3 selected features we defined earlier (including our *custom Feature*), the best model was still the *SVM*, but the hyperparameters assumed different values. In this case, the worse model was instead the *Decision Tree*.

	model	best_score	best_params
0	SVM	0.997138	{'C': 1, 'gamma': 0.01, 'kernel': 'linear'}
1	Random_Forest	0.997138	{'max_depth': 5, 'n_estimators': 20}
2	Logistic_Regression	0.997138	{'C': 1}
3	KNN	0.997138	{'metric': 'minkowski', 'n_neighbors': 11}
4	Decision_Tree	0.996852	{'max_depth': 5}

Table 2.2: Table of scores for selected-features models

Figure 2.3: Best selected-features model confusion matrix when we consider our *custom feature*

Validation

The accuracy we got using only 3 selected features is dramatically high, equal to 0.9977116704805492.

2.6.3 Selected-features model - excluding our *custom feature*

For completeness, using the 3 selected features excluding our *custom feature*, this time the best model turned out to be the *KNN*, while the *Support Vector Machine* loses almost 0.1 points on score. In this case, the worse model was instead the *Decision Tree*. The accuracy we got is equal to 0.9267734553775744. Our model loses near to the 7% of accuracy with the respect to the previous case. Plotting the *confusion matrix* we could see that, while "Distortion" class is still well distinguished, "NoFx" is easily confused with "Tremolo".




Figure 2.4: Best selected-features model confusion matrix when we do not consider our *custom feature*




Figure 2.5: Features plot with the 3 k-best selected features




Figure 2.6: Scatter plot with the 3 k-best selected features

3 | Results And Discussions

3.1 Overall Conclusions On Accuracy Of Our Model

Our model is able to distinguish the 3 classes of sound with high precision. Best results are obtained in distinguishing monophonic electric guitar recordings. For what concern electric bass guitar, a slight decrease in accuracy occurs, but precision remains high. If we do not take into account our *custom feature*, the 3 selected features are really effective in distinguishing "Distortion", but do not provide a good classification of "NoFx" sounds, which are often mistaken for "Tremolo" sounds. To overcome this inaccuracy, we tried to define a *custom feature* that was expected to be selected among the 3 selected features. A proper definition of our *custom feature* led us to well discriminate the two classes and determined a dramatic increase in accuracy. We didn't take into account polyphonic sounds.

Bibliography

- [1] M. Stein. Idmt-smt-audio-effects. [Online]. Available: https://www.idmt.fraunhofer.de/en/business_units/m2d/smt/audio_effects.html
- [2] F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, J. Vanderplas, A. Passos, D. Cournapeau, M. Brucher, M. Perrot, and E. Duchesnay, "Scikit-learn: Machine learning in Python," *Journal of Machine Learning Research*, vol. 12, pp. 2825–2830, 2011.
- [3] B. McFee, V. Lostanlen, A. Metsai, M. McVicar, S. Balke, C. Thomé, C. Raffel, F. Zalkow, A. Malek, Dana, K. Lee, O. Nieto, J. Mason, D. Ellis, E. Battenberg, S. Seyfarth, R. Yamamoto, K. Choi, viktorandreevichmorozov, J. Moore, R. Bittner, S. Hidaka, Z. Wei, nullmightybofo, D. Hereñú, F.-R. Stöter, P. Friesch, A. Weiss, M. Vollrath, and T. Kim. (2020, Jul.) librosa/librosa: 0.8.0. [Online]. Available: <https://doi.org/10.5281/zenodo.3955228>
- [4] H. Mujtaba. (2020) Hyperparameter tuning with gridsearchcv. [Online]. Available: <https://www.mygreatlearning.com/blog/gridsearchcv/>

A1 | List of features

List and brief descriptions of features we used:

- ***Tremolo feature***: The main idea is to count the number of local maxima in the envelope, which are due to the oscillation introduced by the "Tremolo effect". In order to do so we compute the RMS. To avoid the detection of spurious local maxima firstly we apply a threshold and a low pass filter to suppress high frequencies. We use a boolean array to count the number of the remaining local maxima and we return a boolean value: 0 if they are at most 2 (which is a value typical for "NoFx" recordings), or 1 if they are more than 2 peaks (which is typical for "Tremolo" recordings).
- ***Spectral Centroid***: The magnitude spectrogram is treated as a distribution over frequency bins, and the centroid is extracted per frame.

$$centroid = \frac{\sum_k S(k) * freq(k)}{\sum_j S(j)}$$

where $S()$ is the magnitude spectrum.

- ***Spectral Roll-off***: The center frequency for a spectrogram bin such that at least 85% (or another selectable percentage) of the energy of the spectrum in the frame is contained in the bins up to that one.
- ***Spectral Flatness***: Spectral flatness (or tonality coefficient) is a measure to quantify how much noise-like a sound is, as opposed to being tone-like.
- ***Zero-Crossing Rate***: Number of times that the audio waveform crosses the zero

axis.

$$ZCR = \frac{1}{2} \sum_{t=1}^{T-1} |signs(t) - signs(t-1)| \frac{F_s}{N}$$

- **Mel-scaled Spectrogram:** Compute a mel-scaled spectrogram. The mel scale is a logarithmic scale which allows to represent frequencies in a way close to the perception of pitch by a human listener, according to this formula:

$$Pitch(mel) = 1127.0148 * \log(1 + \frac{f}{700})$$

- **Mel-Frequency Cepstral Coefficients:** Those coefficients are computed starting from a series of triangular filter equally spaced on a mel-scale just discussed. At the output of each filter is applied the logarithm, and a Discrete Cosine Transform (DCT).
- **Spectral Contrast:** Each frame of a spectrogram is divided into sub-bands. For each sub-band, the energy contrast is estimated by comparing the mean energy in the top quantile (peak energy) to that of the bottom quantile (valley energy). High contrast values generally correspond to clear, narrow-band signals, while low contrast values correspond to broad-band noise.
- **Root-Mean-Square:** Compute root-mean-square (RMS) value for each frame.

A2 | List of models

List and brief descriptions of models we compared.

- ***Support Vector Machines***: An SVM training algorithm builds a model that assigns new examples to one category or the other, making it a non-probabilistic binary linear classifier. SVM maps training examples to points in space so as to maximise the width of the gap between the two categories. New examples are then mapped into that same space and predicted to belong to a category based on which side of the gap they fall.
- ***Random Forest***: Random forests operates by constructing a multitude of decision trees at training time and outputting the class that is the mode of the classes or mean/average prediction of the individual trees.
- ***Logistic Regression***: Logistic regression is a statistical model that uses a logistic function to model a binary dependent variable. In regression analysis, logistic regression is estimating the parameters of a logistic model. Outputs with more than two values are modeled by multinomial logistic regression.
- ***K-Nearest Neighbors***: In KNN an object is classified by a plurality vote of its neighbors, with the object being assigned to the class most common among its k nearest neighbors.
- ***Decision Tree***: In machine learning a decision tree is a predictive model where each internal node represents a variable while leaves contains the value for a certain variable.